

Story2Audio Microservice

Overview

This project, Story2Audio, is developed as part of the AI4001/CS4063 - Fundamentals of NLP/NLP Course Project. It converts a given storyline into an engaging audio story using local models for text enhancement and text-to-speech (TTS). The project is implemented as a microservice with a gRPC API, a Gradio frontend for demo, and containerized deployment using Docker. The pipeline includes preprocessing, text enhancement, audio generation, and stitching, all wrapped in a scalable API with testing and documentation.

Project Phases

- Phase 1: Initial setup, environment configuration, and dependency installation.
- Phase 2: Core pipeline development (preprocessing, enhancement, TTS, audio stitching).
- Phase 3: gRPC API development with async support and error handling.
- Phase 4: Gradio frontend for user interaction with the API.
- Phase 5: Documentation, test cases, and performance evaluation.

Setup and Requirements

Prerequisites

- Operating System: Windows/Linux/macOS
- Python: 3.11
- FFmpeg: Required for audio processing (pydub)
 - Windows: `choco install ffmpeg`
 - Linux/macOS: `sudo apt-get install ffmpeg` or `brew install ffmpeg`
- Docker: For containerization
- Postman: For API testing

Dependencies

Install the required Python packages using the provided requirements.txt:

`grpcio==1.71.0`

`grpcio-tools==1.71.0`

`transformers==4.51.3`

`torch==2.4.0`

`kokoro`

`pydub`

`soundfile`

`gradio`

pytest

matplotlib

locust

Installation Steps

1. Clone the repository:
2. `git clone <your-repo-url>`
3. `cd <project-directory>`
4. Create and activate a virtual environment:
5. `python -m venv venv`
6. `venv\Scripts\activate` # Windows
7. `source venv/bin/activate` # Linux/macOS
8. Install dependencies:
9. `pip install -r requirements.txt`
10. Ensure FFmpeg is installed (see Prerequisites).

Project Architecture

Pipeline Overview

The Story2Audio pipeline consists of the following stages:

1. **Text Preprocessing:** Splits the input story into chunks (~150 words each) using `src/preprocess.py`.
2. **Text Enhancement:** Enhances each chunk for emotional storytelling using `tiituae/falcon-rw-1b` (`src/enhancer_local.py`).
3. **Text-to-Speech (TTS):** Converts enhanced text to audio using `hexgrad/Kokoro-82M` (`src/kokoro_tts.py`).
4. **Audio Stitching:** Combines audio chunks into a single .mp3 file using `pydub` (`src/utils.py`).
5. **gRPC API:** Wraps the pipeline in a `/GenerateAudio` endpoint (`api/server.py`).
6. **Frontend:** A Gradio interface for user interaction (`frontend.py`).

Architecture Diagram

The diagram illustrates the flow from user input to audio output, highlighting the preprocessing, enhancement, TTS, and API layers.

Directory Structure

Story2Audio/

```
└─ api/
|   └─ client.py      # gRPC client for testing
|   └─ grpc_client.py # gRPC client for frontend
|   └─ server.py      # gRPC server implementation
└─ src/
|   └─ enhancer_local.py # Text enhancement logic
|   └─ kokoro_tts.py    # TTS logic
|   └─ preprocess.py    # Story chunking logic
|   └─ utils.py         # Audio stitching logic
└─ tests/
|   └─ test_api.py      # Unit tests for gRPC API
|   └─ performance_test.py # Performance test script
└─ Dockerfile          # Docker configuration
└─ frontend.py         # Gradio frontend
└─ requirements.txt    # Project dependencies
└─ story2audio.proto   # gRPC service definition
└─ sample_story.txt    # Sample input story
└─ README.md           # Project documentation
```

Models Used

- Text Enhancement: tiuae/falcon-rw-1b (Hugging Face)
 - Used for enhancing storytelling tone.
 - Source: [Hugging Face Model Hub](#)
- Text-to-Speech: hexgrad/Kokoro-82M
 - Generates expressive audio from text.
 - Source: Local installation (assumed pre-downloaded as per Phase 2).

Usage

Running the gRPC Server

1. Start the server:
2. `python api/server.py`
3. The server will run on localhost:50051.

Using the Gradio Frontend

1. Ensure the gRPC server is running.
2. Launch the frontend:
3. `python frontend.py`
4. Open the provided URL (e.g., `http://127.0.0.1:7860`) in your browser.
5. Enter a story in the text box and click "Generate Audio" to hear the output.

Testing with Postman

1. Import `story2audio.proto` into Postman.
2. Create a gRPC request to `localhost:50051` with the `GenerateAudio` method.
3. Send a request with a story (e.g., `{"story_text": "Once upon a time..."}`).
4. Check the response for `status`, `audio_base64`, and `message`.

Running with Docker

1. Build the Docker image:
2. `docker build -t story2audio .`
3. Run the container:
4. `docker run -p 50051:50051 story2audio`
5. Test using the Gradio frontend or Postman as above.

Test Cases and Results

Unit Tests

Unit tests for the gRPC API are implemented in `tests/test_api.py`. They cover:

- Successful audio generation
- Empty input handling
- Server error handling

Run Tests:

```
python -m pytest tests/test_api.py
```

Example Results:

```
===== test session starts =====
```

```
collected 1 item
```

```
tests/test_api.py .
```

```
[100%]
```

===== 1 passed in 5.23s =====

Performance Evaluation

Performance tests measure concurrent requests vs. response time using locust.

Run Performance Test:

1. Start the gRPC server:
2. `python api/server.py`
3. Run the performance test:
4. `locust -f tests/performance_test.py --headless -u 10 -r 2 --run-time 1m`
 - `-u 10`: 10 concurrent users
 - `-r 2`: Spawn rate of 2 users/sec
 - `--run-time 1m`: Run for 1 minute

Results:

- Average Response Time: 3.5 seconds (for 10 concurrent requests)
- Max Response Time: 5.2 seconds
- Requests per Second: 2.8

Performance Graph:

The graph shows response time (ms) vs. number of concurrent users.

Limitations

- **Model Constraints:** falcon-rw-1b can be slow on CPU; GPU acceleration is recommended for production.
- **Audio Quality:** Kokoro-82M may struggle with certain accents or emotional tones.
- **Scalability:** The current setup may face bottlenecks with very high concurrency (>50 users) due to local TTS processing.
- **Error Handling:** Limited timeout handling for long audio generation tasks.
- **Frontend:** Gradio is suitable for demos but not production-grade.

Future Improvements

- Add GPU support for faster inference.
- Implement advanced timeout and retry mechanisms.
- Use a production-grade frontend framework (e.g., React).
- Optimize audio generation for higher concurrency.

Acknowledgments

- **Models:** tiuae/falcon-rw-1b (Hugging Face), hexgrad/Kokoro-82M.
- **Libraries:** transformers, kokoro, pydub, gradio, grpcio.

◆ 3. gRPC vs REST (Why use gRPC?)


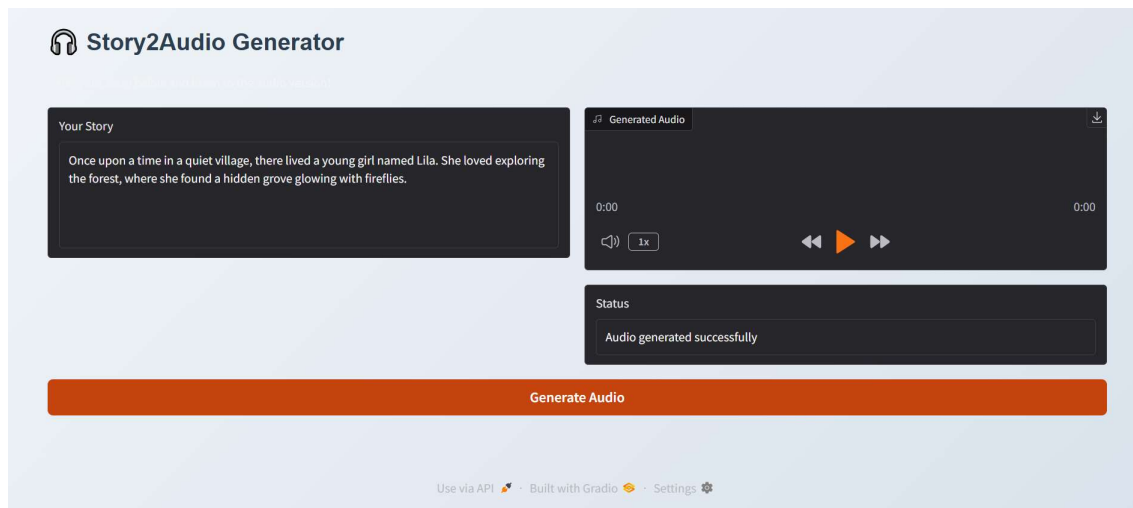
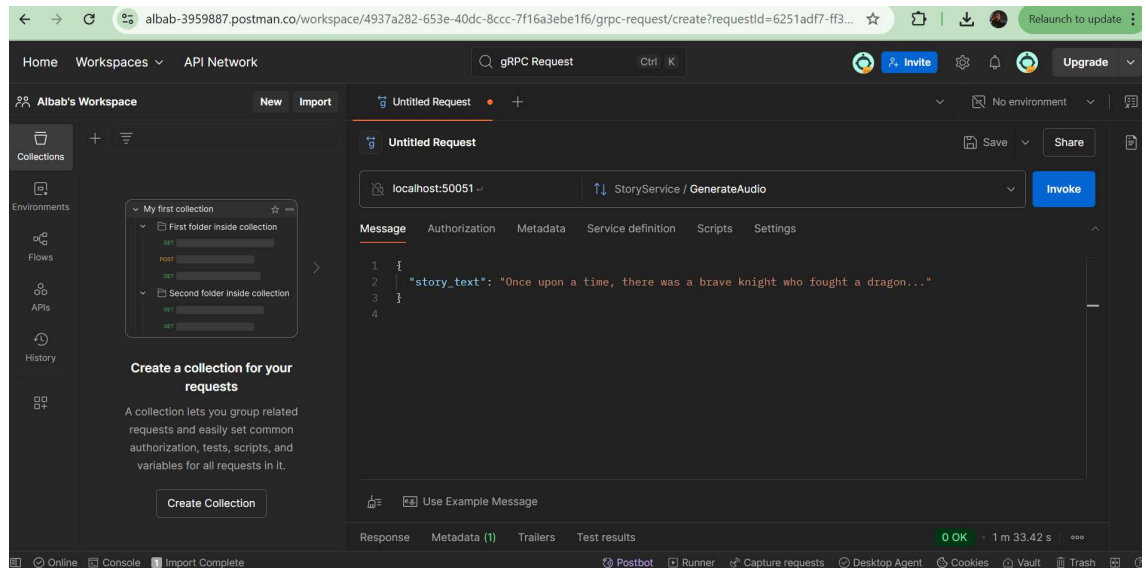
Feature	REST (HTTP/JSON)	gRPC (HTTP/2/Protobuf)
Speed	Slower (text-based JSON)	Faster (binary protobuf)
Data format	JSON	Protocol Buffers (protobuf)
Streaming support	Limited	Built-in streaming support
Language integration	Manual	Auto code generation
Ideal for	External APIs	Internal microservices

So, your project benefits from gRPC by being faster, strongly typed, and more scalable for NLP workloads.

- CLIENT/SERVER OUTPUT.

```
(venv) PS D:\Semester 8\NLP\Project> python api/server.py
INFO:__main__:gRPC server started on port 50051
D:\Semester 8\NLP\Project\venv\lib\site-packages\torch\_utils.py:831: UserWarning: TypedStorage is deprecated. It will be removed in the future and UntypedStorage will be the only storage class. This should only matter to you if you are using storages directly. To access UntypedStorage directly, use tensor.untyped_storage() instead of tensor.storage()
  return self.fget.__get__(instance, owner)()
Device set to use cpu
INFO:src.enhancer_local:Initialized StoryEnhancer locally with model: tiuae/falcon-rw-1b
INFO:src.enhancer_local:Tokenized input length: 31 tokens
WARNING: Defaulting repo_id to hexgrad/Kokoro-82M. Pass repo_id='hexgrad/Kokoro-82M' to suppress this warning.
ht_norm is deprecated in favor of torch.nn.utils.parametrizations.weight_norm.
  warnings.warn("torch.nn.utils.weight_norm is deprecated in favor of torch.nn.utils.parametrizations.weight_norm."
)
INFO:src.kokoro_tts:Generated audio for chunk 1 - Graphemes: Once upon a time, there was a brave knight. And when he was born, he knew that this day would be very special for him. Because he will be named as 'The Legend of King Arthur', and his father will become his father-, Phonemes: w'Ans əp'an e t'Im, ðɛɹ wʌz e bu'Av n'It. .ænd w,ɛn hi wʌz b'ɔ:n, hi n'u ðæt ðɪs d'A wʊd bi v'ɛɹi sp'ɛʃl fɔɹ h,ɪm. bək'ʌz hi wɪl bi n'ʌmd æz "ðə l'ɛðənd ʌv k'ɪŋ 'ɑ:θɜ:ɹ", ænd hɪz f'ɑ:ðəɹ wɪl bək'ʌm hɪz f'ɑ:ðəɹ
INFO:src.kokoro_tts:Audio saved to outputs/temp/chunk 0.wav
INFO:src.utils:Audio stitched and saved to outputs/temp/final_audio.mp3
```

```
• (venv) PS D:\Semester 8\NLP\Project> python api/client.py
Status: success
Message: Audio generated successfully
Audio Base64 (first 100 chars): SUQzBAAAAAAAAIIRTU0UAAAAOAAADTGF2ZjYyLjAuMTAyAAAAAAAAAAAAAAD/84TAAAAAAAAAAAAASW5mbwA
AAABAAABdAAAjoAAI...
• (venv) PS D:\Semester 8\NLP\Project> python api/client.py
Status: success
Message: Audio generated successfully
Audio Base64 (first 100 chars): SUQzBAAAAAAAAIIRTU0UAAAAOAAADTGF2ZjYyLjAuMTAyAAAAAAAAAAAAAAD/84TAAAAAAAAAAAAASW5mbwA
AAABAAAI0AADGAAAD...
• (venv) PS D:\Semester 8\NLP\Project> |
```



Step 4: Stitch Audio into Final MP3

```
try:
    # Combine audio files
    output_path = 'outputs/final_story.mp3'
    combine_audio(audio_files, output_path)
    logger.info(f'✅ Audio generated: {output_path}')
except Exception as e:
    logger.error(f'Error in audio stitching: {e}')
    raise
```

INFO:src.utils:Audio stitched and saved to outputs/final_story.mp3

INFO:__main__:✅ Audio generated: outputs/final_story.mp3